

Reference: [ML System Design by ByteByteGo](#).




Figure 1: Components of a production ready ML system.

Introduction and Overview

ML system design interviews expect us to answer open-ended questions. There is no single correct answer. We are evaluated on our thought process, our in-depth understanding of various ML topics, our ability to design an end-to-end system, and our design choices based on the trade-offs of various options.

Unstructured answers make the flow difficult to follow and derail open-ended discussions. It is important to follow a framework:




Figure 2: ML system design steps.

- Clarify the objective to help shape primary requirements. Transform the objective/ requirements into a tangible ML task.
- Data preparation is often overlooked. Don't be an idiot. Think through what will be fed to the model, both in terms of X and Y .

- Model development and evaluation must be an continuously iterative process. To **MeasureEverything** is important; ideally evaluation (and non-ML baselines) must be decided in the early stages of model development.
- Continuous monitoring after serving is essential to ensure that the objectives are being met and that the data/ concept does not drift.

Clarify Requirements

Ask clarifying questions; asking questions exhibits attention to detail and intention. Attempt to understand the exact requirements and what the user (the interviewer) cares about:

- **Business objective:** Increase revenue? Increase #users? Increase engagement per user?
- **Features the system needs to support:** Move towards tangible features. Clarify if the features are essential or good-to-haves. Discuss feasibility while keeping a realistic ML system in mind. For example, a recommendation system might allow users to “like” or “dislike” recommendations, as those interactions could be used to label training data.
- **Data:** Data sources? How large is a datum and the dataset? Is it structured or unstructured? Do we have labels and is there a possibility for soft labels?
- **Constraints:** How much computing power is available (during inference; during learning we assume ∞ resources)? Is it a cloud-based system, or should the system work on a device? Is the model expected to improve automatically over time?
- **Scale of the system:** How many users do we have? How many items, such as videos, are we dealing with? What’s the rate of growth of these metrics?
- **Performance:** How fast must inference be? Is a real-time solution expected? Does accuracy have more priority or latency?
- **FairML:** Do Fairness, Accountability, Transparency, and Explainability matter?

Frame the Problem as an ML Task

Effective problem framing is essential for translating a business objective into a clear ML task. This involves (1) ensuring that ML is necessary (though an ML system design interview probably requires an ML system...), (2) defining a precise ML objective that models can address (e.g., maximizing click-through

rate instead of merely “increasing sales”), (3) specifying the system’s inputs and outputs as a blackbox — especially when multiple models are involved, and (4) choosing the appropriate ML paradigm (e.g., supervised learning for most cases, with further distinctions between binary/ multiclass classification or regression), ensuring the task aligns with the nature of available data and the intended outcome.

Define the ML Objective Possible business objective: increase sales by 20%. But, we cannot optimize over sales directly. For an ML system to solve a task, we must translate the business objective into a well-defined ML objective (say, to improve recommendations so users buy more items). We could also increase sales by increasing the #users by 20%, but that might require moving away from improving recommendations towards, say, a better ad-serving system. Increasing sales can only be measured on the field and that is not a good ML objective for learning and evaluating the ML system.

A *good ML objective* is one that ML models can solve and be evaluated upon.

Specify the System’s I/O As a black-box what are X and Y? In some cases, the system may be complex and we may need to specify the I/O of each component ML model. There may be more than one way to specify the I/O and these decisions matter for how we make downstream decisions.




Figure 3: Different ways to specify the model’s I/O.